

AI Chatbot & Assistant

Custom AI assistants that understand your business and delight your customers

• Copenhagen, Denmark

9 LLM Providers

Multi-model Fallback

RAG-powered

View the full interactive portfolio with live demos and animations

[View Interactive Portfolio](#)

What I Build

Click any card to see the full breakdown.

Customer Support Bot

24/7 automated support, FAQ handling, ticket routing, human handoff

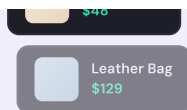
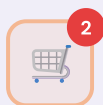


Section 4.2

HR Policy

Knowledge Base Assistant

RAG-powered search over your docs, SOPs, wikis with citations



E-commerce Concierge

Product recommendations, order tracking, returns, upselling

Engaged by bot

412

Meeting booked

Lead Qualification Bot

Qualify visitors, book meetings, capture info, CRM sync

ask-hr-bot



Sarah

How many vacation days do I have left?

Internal Team Assistant

Slack/Teams bot for HR, IT, onboarding, policy lookups



Voice & Multimodal

Voice-enabled assistants, image recognition, document processing

Live System: ElektraOS

Production AI assistant with cross-session memory and multi-provider intelligence.

ElektraOS

AI CHAT SYSTEM

```
$ elektra chat --mode=rag
INFO: mem0 memory loaded
INFO: vault RAG indexed
• Elektra online
```

ElektraOS is a live AI assistant system featuring the Elektra chat engine with cross-session memory (mem0), vault-based RAG semantic search, multi-provider LLM fallback chains across 9 providers, and 624+ API endpoints. It remembers past conversations, searches internal knowledge bases, and gracefully falls back between models for reliability.

0

LLM Providers

0

Cross-session Memory

0

RAG Search

0

Endpoints

Tech Stack

Groq

Anthropic

OpenAI

Ollama

mem0

ChromaDB

FastAPI

Python

WebSockets

LangChain

DSPy

OpenLIT

Docker

How It Works

From idea to live chatbot in 5 clear steps.



1. Discovery

Understand your business, users, and use cases. Identify the highest-impact conversations to automate first.



2. Design Conversation Flows

Map intents, responses, fallback paths, and human handoff rules. Define the bot's personality and knowledge scope.



3. Build & Train

Create a custom knowledge base, build the RAG pipeline, tune the model personality, and test with real conversation data.



4. Integrate

Connect to your website, Slack, WhatsApp, CRM, or helpdesk. Set up monitoring, analytics, and alert hooks.



5. Launch & Optimize

Deploy to production, monitor live conversations, and improve accuracy weekly based on real user interactions.



AI Quality & Investment

The quality of AI responses scales directly with the language model tier used. Entry-level models (GPT-3.5, Groq Llama) deliver solid results for FAQ and simple queries. Premium models (GPT-4, Claude Opus, Gemini Pro) unlock deeper reasoning, better context handling, and more natural conversations – but come at higher API costs. I always recommend the right model for your use case and budget, with transparent pricing on API costs passed through at cost.

Pricing

Transparent pricing. All tiers include testing, documentation, and support.

FAQ Bot

\$200 USD

- Simple Q&A bot
- Up to 50 FAQs
- Website chat widget
- Basic analytics
- 5-day delivery

POPULAR

Smart Assistant

\$800 USD

- RAG-powered knowledge base
- Custom training data
- Multi-channel deployment
- CRM integration
- Human handoff
- 10-day delivery

Enterprise AI

\$2000+ USD

- Multi-agent system
- Voice support
- Custom LLM fine-tuning
- SLA guarantee
- Dedicated support
- Custom timeline

Ready to give your customers 24/7 AI-powered support?

Stop losing leads to slow response times. Let's build an AI assistant that works around the clock.

[Build My Chatbot →](#)

or email: yossra.benzad@gmail.com